

Reinforcement Learning



UPC DLAI Group8: Adrian, Bruno, Gianmarco and Jordi.

Agenda

- Introduction
- Deep Stochastic Policy Gradient Agent
- Deep Stochastic Policy Gradient Agent Experiments
- Deep Q Network Agent
- Conclusions

Introduction

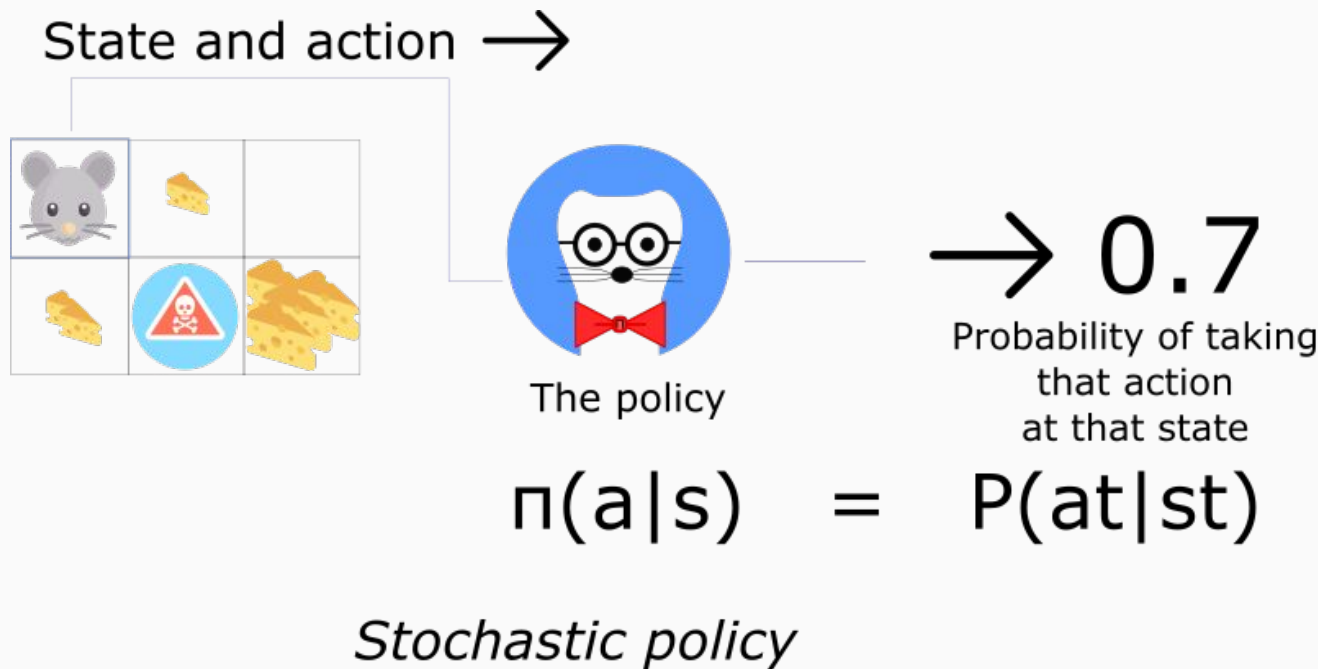
Reinforcement Learning



“... Reinforcement learning is learning what to do—how to map situations to actions—so as to maximize a numerical reward signal. The learner is not told which actions to take, but instead must discover which actions yield the most reward by trying them...” Sutton, R. S.

Deep Stochastic Policy Gradient Agent I

- **Basics:**
 - **Stochastic:** Our policy outputs a probability distribution.
 - **Deep:** It is modelled by a deep neural network
 - **Policy:** Maps an state s to a probability distribution among actions
 - **Agent:** Takes actions by sampling from the policy distribution



Deep Stochastic Policy Gradient Agent II

Learning the policy:

- Measure how good is our policy using an **Objective Function J**
- Compute gradients:

$$J_1(\theta) = V_{\pi_\theta}(s_1) = E_{\pi_\theta}[v_1] = \underbrace{\sum_{s \in S} d(s)}_{\text{State distribution}} \underbrace{\sum_{a \in A} \pi_\theta(s, a) R_s^a}_{\text{Action distribution}}$$

$$\nabla_\theta J(\theta) = E_\pi[\underbrace{\nabla_\theta(\log \pi(\tau|\theta))}_{\text{Policy function}} \underbrace{R(\tau)}_{\text{Score function}}]$$

- Update policy parameters
Monte Carlo update



function REINFORCE

Initialise θ arbitrarily

for each episode $\{s_1, a_1, r_2, \dots, s_{T-1}, a_{T-1}, r_T\} \sim \pi_\theta$ **do**

for $t = 1$ to $T - 1$ **do**

$\theta \leftarrow \theta + \alpha \nabla_\theta \log \pi_\theta(s_t, a_t) v_t$

end for

end for

return θ

end function

Deep Stochastic Policy Gradient Agent III

Tanh vs ReLU activations

Param:

learning_rate \rightarrow 0.001


optimizer \rightarrow adam

batch_size = 32

n_experiments= 10

Blue: FCN using tanh activations

Red: FCN using ReLU activations



Deep Stochastic Policy Gradient Agent III

Entropy Normalization:

Encouraging our agent to explore the environment.

Param:


learning_rate \rightarrow 0.001

optimizer \rightarrow adam

batch_size = 32

n_experiments= 10

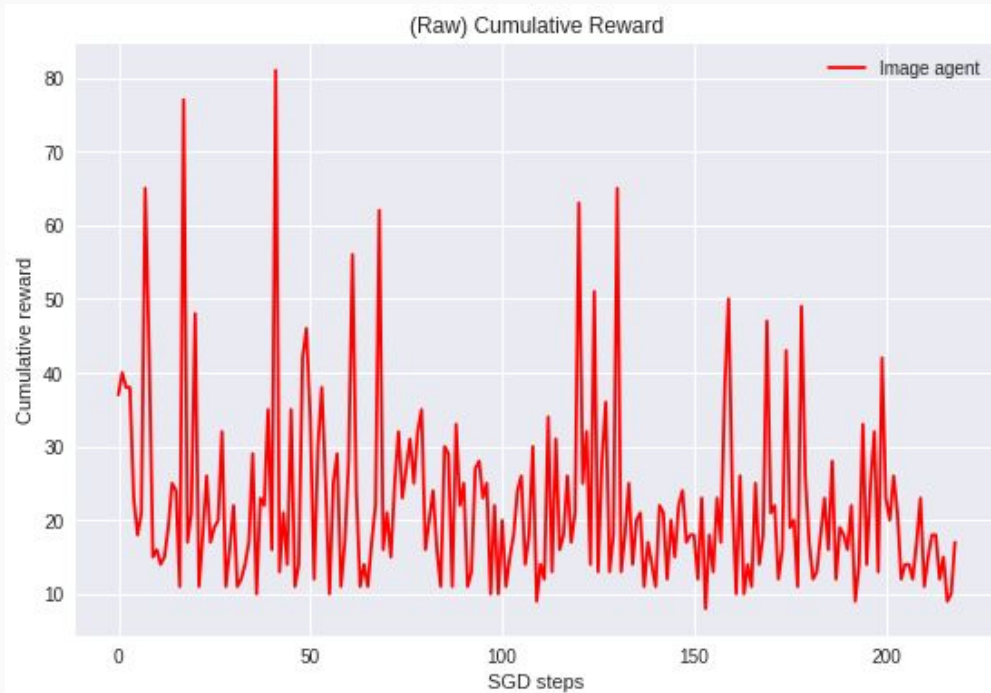
Based on the publication Understanding the impact of entropy on policy optimization, Nov 2018



Deep Stochastic Policy Gradient Agent III


State as RGB frame

- Learning the policy from images is much more challenging.
- New architecture: 3 convolution layers + 2 FC + softmax activation.
- Poor results



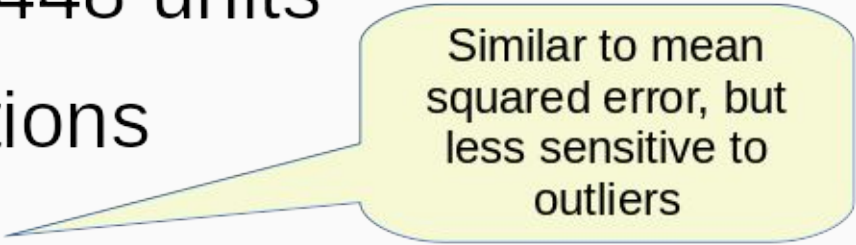
Reinforcement learning with DQN

- State
 - Current frame from the environment
 - Cropped
 - Previous frame subtracted
- Replay memory
 - Memorize the previous experience
 - Randomly samples



The Network

- 3 convolutional layers with batch normalization
 - Size 16-32-32
- One linear layer with 448 units
- ReLU activation functions
- Huber loss function
- RMSprop optimizer



Similar to mean squared error, but less sensitive to outliers

Training results

- Our network did not converge
- Changing some Hyper-parameter did not help
- Perhaps it is possible to tune them better



Conclusions

- We implemented two different approaches with different activation functions
- Real environments like an RGB image state are much more complex.
- Encouraging exploration becomes essential in complex environments.
- Reinforcement learning is not easy to converge

Thank you for your attention!

Reinforcement Learning

UPC DLAI Group8: Adrian, Bruno, Gianmarco and Jordi.